

Amazon Fine Food Reviews Analysis

Data Source: <https://www.kaggle.com/snap/amazon-fine-food-reviews>

EDA: <https://nycdatascience.com/blog/student-works/amazon-fine-foods-visualization/>

The Amazon Fine Food Reviews dataset consists of reviews of fine foods from Amazon.

Number of reviews: 568,454

Number of users: 256,059

Number of products: 74,258

Timespan: Oct 1999 - Oct 2012

Number of Attributes/Columns in data: 10

Attribute Information:

1. Id
2. ProductId - unique identifier for the product
3. UserId - unique identifier for the user
4. ProfileName
5. HelpfulnessNumerator - number of users who found the review helpful
6. HelpfulnessDenominator - number of users who indicated whether they found the review helpful or not
7. Score - rating between 1 and 5
8. Time - timestamp for the review
9. Summary - brief summary of the review
10. Text - text of the review

Objective:

Given a review, determine whether the review is positive (rating of 4 or 5) or negative (rating of 1 or 2).

[Q] How to determine if a review is positive or negative?

[Ans] We could use Score/Rating. A rating of 4 or 5 can be considered as a positive review. A rating of 1 or 2 can be considered as negative one. A review of rating 3 is considered neutral and such reviews are ignored from our analysis. This is an approximate and proxy way of determining the polarity (positivity/negativity) of a review.

[1]. Reading Data

[1.1] Loading the data

The dataset is available in two forms

1. .csv file
2. SQLite Database

In order to load the data, We have used the SQLITE dataset as it is easier to query the data and visualise the data efficiently.

Here as we only want to get the global sentiment of the recommendations (positive or negative), we will purposefully ignore all Scores equal to 3. If the score is above 3, then the recommendation will be set to "positive". Otherwise, it will be set to "negative".

In [1]:

```
import sqlite3
import pandas as pd
import numpy as np
import nltk
import string
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.feature_extraction.text import TfidfTransformer
from sklearn.feature_extraction.text import TfidfVectorizer
```

```

from sklearn.model_selection import train_test_split
from sklearn.metrics import accuracy_score
from sklearn.model_selection import cross_val_score
from collections import Counter
from sklearn.preprocessing import StandardScaler

from sklearn.feature_extraction.text import CountVectorizer
from sklearn.metrics import confusion_matrix
from sklearn import metrics
from sklearn.metrics import roc_curve, auc

import re
# Tutorial about Python regular expressions: https://pymotw.com/2/re/
import string
from nltk.corpus import stopwords
from nltk.stem import SnowballStemmer
from nltk.stem.wordnet import WordNetLemmatizer

import pickle

```

In [1]:

```

data = pd.read_csv('Reviews.csv')
print (data.head(2))
print(data.shape)

```

```

   Id  ProductId      UserId ProfileName  HelpfulnessNumerator  \
0   1  B001E4KFG0  A3SGXH7AUHU8GW  delmartian                    1
1   2  B00813GRG4  A1D87F6ZCVE5NK      dll pa                    0

   HelpfulnessDenominator  Score      Time      Summary  \
0                        1      5  1303862400  Good Quality Dog Food
1                        0      1  1346976000    Not as Advertised

                                     Text
0  I have bought several of the Vitality canned d...
1  Product arrived labeled as Jumbo Salted Peanut...
(568454, 10)

```

[2] Data Cleaning: Deduplication and Nan features

It is observed (as shown in the table below) that the reviews data had many duplicate entries. Hence it was necessary to remove duplicates in order to get unbiased results for the analysis of the data. Following is an example:

In [2]:

```

#checking for Nan values in data. True indicates Nan values are present along the columns
data.isnull().any()

```

Out[2]:

```

Id                False
ProductId          False
UserId            False
ProfileName        True
HelpfulnessNumerator  False
HelpfulnessDenominator  False
Score              False
Time              False
Summary            True
Text               False
dtype: bool

```

In [3]:

```

# checking for Nan values along 'profilename' column
#data[data['ProfileName'].isnull()].head(2)

```

In [4]:

```

# checking for Nan values along 'summary' column

```

```
# checking for nan values along 'summary' column
#data[data['Summary'].isnull()]
```

In [5]:

```
#Dropping Nan values
data = data.dropna()
```

In [6]:

```
#printing shape of data after dropping Nan values
print (data.shape)
```

(568411, 10)

In [7]:

```
#Review score should lie between 1 to 5
#Returns True if all the scores lie between 1 to 5(inclusive)
list1 = data['Score'].map(lambda x: True if x in [1,2,3,4,5] else False)
list1.all()
```

Out[7]:

True

In [8]:

```
filtered_data = data.loc[data['Score']!=3].copy()
print (filtered_data.head(2))
print (filtered_data.shape)
```

| | Id | ProductId | UserId | ProfileName | HelpfulnessNumerator | \ |
|---|----|------------|----------------|-------------|----------------------|---|
| 0 | 1 | B001E4KFG0 | A3SGXH7AUHU8GW | delmartian | 1 | |
| 1 | 2 | B00813GRG4 | A1D87F6ZCVE5NK | d11 pa | 0 | |

| | HelpfulnessDenominator | Score | Time | Summary | \ |
|---|------------------------|-------|------------|-----------------------|---|
| 0 | 1 | 5 | 1303862400 | Good Quality Dog Food | |
| 1 | 0 | 1 | 1346976000 | Not as Advertised | |

| | Text |
|---|---|
| 0 | I have bought several of the Vitality canned d... |
| 1 | Product arrived labeled as Jumbo Salted Peanut... |

(525773, 10)

In [9]:

```
#mapping positive(>3) and negative(<3) reviews based on scores of the data.
import pandas as pd
pos_negative = filtered_data['Score'].map(lambda x: 1 if int (x)>3 else 0)
filtered_data['Score'] = pos_negative
print ('shape of filtered_data')
print (filtered_data.shape)
#print (filtered_data.head())
```

shape of filtered_data
(525773, 10)

In [10]:

```
#arranging data with increasing productid
sorted_data = filtered_data.sort_values('ProductId',axis=0,ascending=True,inplace=False,kind='quick
sort')
```

In [11]:

```
#finding the duplicates in our data
#If the same person gives for the same product at the same time we call it as suplicates
#sorted_data.loc[sorted_data.duplicated(["UserId","ProfileName","Time","Text"],keep = False):1
```

```
#sorted_data=sorted_data.duplicated(['UserId','ProfileName','Time','Text'],keep='first',inplace=False,/)


```

In [12]:

```
#counting number of duplicates present in our data
sorted_data.duplicated(["UserId","ProfileName","Time","Text"]).sum()
```

Out[12]:

161612

In [13]:

```
#dropping all duplicates keeping the first one
final = sorted_data.drop_duplicates(subset={"UserId","ProfileName","Time","Text"},keep='first',inplace = False)
final.shape
```

Out[13]:

(364161, 10)

In [14]:

```
#Checking to see how much % of data still remains
(final['Id'].size*1.0)/(sorted_data['Id'].size*1.0)*100
```

Out[14]:

69.26201992114468

In [15]:

```
#helpfulness numerator denotes number of people who found the review helpful
#helpfulness denominator denotes number of people who indicated whether or not the review helpful
#so, helpfulness numerator should be less than denominator
final = final[final.HelpfulnessNumerator<=final.HelpfulnessDenominator]
```

In [16]:

```
#final shape of data after preprocessing
final.shape
```

Out[16]:

(364159, 10)

In [17]:

```
final['Score'].value_counts()
```

Out[17]:

```
1    307054
0     57105
Name: Score, dtype: int64
```

In [18]:

```
final.shape
```

Out[18]:

(364159, 10)

In [19]:

```
#arranging data with increasing time
final_data = final.sort_values('Time',axis=0,ascending=True,inplace=False,kind='quicksort')
```

[3] Preprocessing

[3.1]. Preprocessing Review Text

Now that we have finished deduplication our data requires some preprocessing before we go on further with analysis and making the prediction model.

Hence in the Preprocessing phase we do the following in the order below:-

1. Begin by removing the html tags
2. Remove any punctuations or limited set of special characters like , or . or # etc.
3. Check if the word is made up of english letters and is not alpha-numeric
4. Check to see if the length of the word is greater than 2 (as it was researched that there is no adjective in 2-letters)
5. Convert the word to lowercase
6. Remove Stopwords
7. Finally Snowball Stemming the word (it was observed to be better than Porter Stemming)

After which we collect the words used to describe positive and negative reviews

In [120]:

```
#final data = final_data.iloc[0:10000,:].copy()
```

In [20]:

```
stop = set(stopwords.words('english')) #set of stopwords
sno = nltk.stem.SnowballStemmer('english') #initialising the snowball stemmer
not_words = re.findall(r'\w*n[\'|o]t',str(stop)) #finding NOT words in stop words
not_words.append('\n\t')
not_words.append('no')
print(not_words)
stop_words = stop - set(not_words) #removing NOT words from stop words
# https://stackoverflow.com/a/47091490/4084039
import re

def decontracted(phrase):
    # specific
    phrase = re.sub(r'\w*n[\'|o]t', "not", phrase)
    # general
    phrase = re.sub(r"\'re", " are", phrase)
    phrase = re.sub(r"\'s", " is", phrase)
    phrase = re.sub(r"\'d", " would", phrase)
    phrase = re.sub(r"\'ll", " will", phrase)
    phrase = re.sub(r"\'t", " not", phrase)
    phrase = re.sub(r"\'ve", " have", phrase)
    phrase = re.sub(r"\'m", " am", phrase)
    return phrase

def cleanhtmlpunc(sentence): #function to clean the word of any html-tags
clean = re.compile('<.*?>')
clean = re.sub(clean, ' ', sentence)
clean = re.sub(r"(http|www)\S+", "", clean)
clean = re.sub(r"\S+com", "",clean)
#clean = re.sub(r"(\w+)\.", "",clean)
clean = re.sub(r"\.", "" ,clean)
cleaned = re.sub(r'[?!+|\'+|"+|#|:]',r' ',clean)
cleantext = re.sub(r'[\.\+,+]+|(+\|\/+]',r' ',cleaned)
return cleantext

#def cleanpunc(sentence): #function to clean the word of any punctuation or special characters
# return cleaned

print('*****')
print(stop_words)
```

```
['couldn't', 'haven't', 'shan't', 'wouldn't', 'mightn't', 'not', 'don't', 'weren't', 'hasn't', 'wo  
n't', 'aren't', 'wasn't', 'isn't', 'hadn't', 'mustn't', 'shouldn't', 'doesn't', 'didn't',  
'needn't', 'n't', 'no']  
*****  
{ 'you're', 'wouldn', 'does', 't', 'i', 'd', 'at', 'them', 'didn', 'a', 'when', 'we', 'should've',  
's', 'that'll', 'but', 'if', 'did', 'up', 'are', 'haven', 'the', 'now', 'on', 'hers', 'our', 'most
```

```
, 'his', 'by', 'below', 'doesn', 'as', 'with', 'aren', 'whom', 'her', 'have', 'until', 'that', 'd
uring', 'will', 'other', 'some', 'of', 'doing', 'then', 'each', 'only', 'off', 'than', 'after', 'i
sn', 'and', 'can', 'where', 'what', 'should', 'against', 'couldn', 'from', 'm', 'an', 'about', 'me
', 're', 'o', 'how', 'ain', 'such', 'few', 'itself', 'to', 'yourselves', 'mustn', 'own',
'themselves', 'myself', 'further', 'needn', 'their', "you'd", 'again', 'hasn', 'your', 'had', 'sho
uldn', 'shan', 'has', 'don', 'for', 'he', 'is', 'under', "you'll", 'just', 'nor', 'ma', 'all', 'my
', "she's", 'him', 'above', 'll', 'or', 'so', 'who', 'more', 'out', 'before', 'down', "it's", 'thr
ough', 'were', 'ourselves', 'here', 'yours', 'being', 'wasn', 'they', 'in', 'theirs', 'was', 'be',
'it', "you've", 'too', 'yourself', 'she', 'between', 'these', 'won', 'why', 'y', 'himself', 'you',
'very', 'am', 'having', 'do', 'this', 'which', 'there', 'into', 'same', 'once', 'over', 'ours', 'b
ecause', 'mightn', 'its', 'while', 've', 'any', 'both', 'weren', 'those', 'hadn', 'been',
'herself'}
```

In [21]:

```
def cleanedtext(reviews):
    str1=' '
    final_string=[]
    s=' '
    for sent in reviews:
        filtered_sentence=[]
        sent=cleanhtmlpunc(decontracted(sent)) # remove HTML tags
        for w in sent.split():
            for cleaned_words in w.split():
                if((cleaned_words.isalpha()) & (len(cleaned_words)>2)):
                    if((cleaned_words.lower() not in stop_words)):
                        s=(sno.stem(cleaned_words.lower())).encode('utf8')
                        filtered_sentence.append(s)
                    else:
                        continue
                else:
                    continue
            str1 = b" ".join(filtered_sentence) #final string of cleaned words for reviews
            final_string.append(str1)
    return final_string
```

In [22]:

```
final_string = cleanedtext(final_data['Text'].values)
```

In [23]:

```
final_data['CleanedText']=final_string #adding a column of CleanedText which displays the data aft
er pre-processing of the review
print (final_data.shape)
final_data.head(2) #below the processed review can be seen in the CleanedText Column
```

(364159, 11)

Out[23]:

| | Id | ProductId | UserId | ProfileName | HelpfulnessNumerator | HelpfulnessDenominator | Score | |
|--------|--------|------------|---------------|-----------------------|----------------------|------------------------|-------|------|
| 150523 | 150524 | 0006641040 | ACITT7DI6IDDL | shari zychinski | 0 | 0 | 1 | 9393 |
| 150500 | 150501 | 0006641040 | AJ46FKXOVC7NR | Nicholas A Mesiano | 2 | 2 | 1 | 9408 |

In [2]:

```
import pickle
with open("final_data.pkl", "rb") as f:
```

```
final_data = pickle.load(f)
```

```
In [3]:
```

```
final_data = final_data[0:5000].copy()
```

4) FEATURIZATION

In featurization we use BOW and TF-IDF, as the values in naive bayes need to be positive and word2vec may contain negative values.

(4)APPLYING AGGLOMERATIVE CLUSTERING

4.1) Avg W2V

```
In [9]:
```

```
import pickle
with open("avg_w2v_train_data.pkl", "rb") as f:
    avg_w2v_train_data = pickle.load(f)
```

```
In [10]:
```

```
avg_w2v_train_data1 = avg_w2v_train_data[0:5000]
```

```
In [11]:
```

```
standardizing = StandardScaler(with_mean = False)
avg_w2v_std_train_data = standardizing.fit_transform(avg_w2v_train_data1)
avg_w2v_std_train_data.shape
```

```
Out[11]:
```

```
(5000, 300)
```

```
In [12]:
```

```
from sklearn.cluster import AgglomerativeClustering
aggclu = AgglomerativeClustering(n_clusters=2)
aggclu.fit(avg_w2v_std_train_data)
```

```
Out[12]:
```

```
AgglomerativeClustering(affinity='euclidean', compute_full_tree='auto',
    connectivity=None, linkage='ward', memory=None, n_clusters=2,
    pooling_func=<function mean at 0x7f61680a5ae8>)
```

```
In [13]:
```

```
final_data['clusters'] = aggclu.labels_
final_data['clusters'].value_counts()
```

```
Out[13]:
```

```
0    3971
1     1029
Name: clusters, dtype: int64
```


```
In [14]:
```

```
final_data['Score'].value_counts()
```

```
Out[14]:
```

In [15]:

*****word cloud for cluster0*****



*****word cloud for cluster0*****



*****word cloud for cluster3*****





```
num of leaves 5000
num of components 1
```

In [26]:

```
from wordcloud import WordCloud
for i in [0,2,5]:
    reviews = ''
    for review in final_data[final_data['clusters']==i]['Cleaned'].values:
        reviews = reviews + ' ' + str(review)
    print ('*****word cloud for
cluster{0}*****'.format(i))
    wordcloud = WordCloud(width = 400, height = 400,
                           background_color='white',max_words=400).generate(reviews)
    plt.figure(figsize = (8, 8), facecolor = None)
    plt.imshow(wordcloud)
    plt.axis("off")
    plt.tight_layout(pad = 0)
    plt.show()
```

```
*****word cloud for cluster0*****
```



*****word cloud for cluster2*****




```
tfidf_w2v_train_data1 = tfidf_w2v_train_data[0:5000]
```

In [29]:

```
standardizing = StandardScaler(with_mean = False)
tfidf_w2v_std_train_data = standardizing.fit_transform(tfidf_w2v_train_data1)
tfidf_w2v_std_train_data.shape
```

Out[29]:

(5000, 300)

In [31]:

```
aggclu = AgglomerativeClustering(n_clusters=2)
aggclu.fit(tfidf_w2v_std_train_data)
```

Out[31]:

```
AgglomerativeClustering(affinity='euclidean', compute_full_tree='auto',
    connectivity=None, linkage='ward', memory=None, n_clusters=2,
    pooling_func=<function mean at 0x7f61680a5ae8>)
```

In [32]:

```
final_data['clusters'] = aggclu.labels_
final_data['clusters'].value_counts()
```

Out[32]:

```
0    4644
1     356
Name: clusters, dtype: int64
```

In [33]:

```
final_data['Score'].value_counts()
```


Out[33]:

```
1    4422
0     578
Name: Score, dtype: int64
```

In [35]:

```
from wordcloud import WordCloud
for i in [0,1]:
    reviews = ''
    for review in final_data[final_data['clusters']==i]['Cleaned'].values:
        reviews = reviews + ' ' + str(review)
    print('*****word cloud for
cluster{0}*****'.format(i))
    wordcloud = WordCloud(width = 400, height = 400,
        background_color = 'white',max_words=400).generate(reviews)
    plt.figure(figsize = (8, 8), facecolor = None)
    plt.imshow(wordcloud)
    plt.axis("off")
    plt.tight_layout(pad = 0)
    plt.show()
```

*****word cloud for cluster0*****





*****word cloud for cluster1*****



In [37]:

```
aggclu = AgglomerativeClustering(n_clusters=12)
aggclu.fit(tfidf_w2v_std_train_data)
```

Out[37]:

```
AgglomerativeClustering(affinity='euclidean', compute_full_tree='auto',
connectivity=None, linkage='ward', memory=None, n_clusters=12,
```

```
pooling_func=<function mean at 0x7f61680a5ae8>
```

In [38]:

```
final_data['clusters'] = aggclu.labels_  
final_data['clusters'].value_counts()
```

Out [38] :

| | |
|----|------|
| 0 | 1249 |
| 7 | 749 |
| 3 | 605 |
| 2 | 544 |
| 8 | 400 |
| 1 | 356 |
| 5 | 277 |
| 6 | 238 |
| 9 | 213 |
| 4 | 211 |
| 10 | 99 |
| 11 | 59 |

Name: clusters, dtype: int64

In [39]:

```
from wordcloud import WordCloud
for i in [0,7,3]:
    reviews = ''
    for review in final_data[final_data['clusters']==i]['Cleaned'].values:
        reviews = reviews + ' ' + str(review)
    print ('*****word cloud for
cluster{0}*****'.format(i))
    wordcloud = WordCloud(width = 400, height = 400,
                           background_color = 'white', max_words=400).generate(reviews)
    plt.figure(figsize = (8, 8), facecolor = None)
    plt.imshow(wordcloud)
    plt.axis("off")
    plt.tight_layout(pad = 0)
    plt.show()
```

*****word cloud for cluster0*****



long sweet work took better

*****word cloud for cluster7*****



*****word cloud for cluster3*****





5) APPLYING DBSCAN

5.1) AVG W2V

In [40]:

```
from sklearn.cluster import DBSCAN
dbscan = DBSCAN(eps = 0.5)
dbscan.fit(avg_w2v_std_train_data)
```

Out[40]:

```
DBSCAN(algorithm='auto', eps=0.5, leaf_size=30, metric='euclidean',
       metric_params=None, min_samples=5, n_jobs=1, p=None)
```

In [41]:

```
final_data['clusters'] = dbscan.labels_  
final_data['clusters'].value_counts()
```

Out[41]:

$$\begin{array}{r} -1 \quad 4995 \\ 0 \quad 5 \end{array}$$

```
Name: clusters, dtype: int64
```

In [42]:

```
from wordcloud import WordCloud
for i in [0,-1]:
    reviews = ''
    for review in final_data[final_data['clusters']==i]['Cleaned'].values:
        reviews = reviews + ' ' + str(review)
    print ('*****word cloud for cluster {0}
*****'.format(i))
    wordcloud = WordCloud(width = 400, height = 400,
                           background_color='white',max_words=500).generate(reviews)
    plt.figure(figsize = (8, 8), facecolor = None)
    plt.imshow(wordcloud)
    plt.axis("off")
    plt.tight_layout(pad = 0)
    plt.show()
```

*****word cloud for cluster 0 *****





*****word cloud for cluster -1 *****



In [43]:

```
dbscan = DBSCAN(eps = 1)
dbscan.fit(avq w2v std train data)
```

Out[43]:

```
DBSCAN(algorithm='auto', eps=1, leaf_size=30, metric='euclidean',
       metric_params=None, min_samples=5, n_jobs=1, p=None)
```

In [44]:

```
final_data['clusters'] = dbscan.labels_  
final_data['clusters'].value counts()
```

Out[44]:

| | |
|----|------|
| -1 | 4990 |
| 1 | 5 |
| 0 | 5 |

Name: clusters, dtype: int64

In [45]:

```
dbscan = DBSCAN(eps = 4)
dbscan.fit(avg_w2v_std_train_data)
final_data['clusters'] = dbscan.labels_
final_data['clusters'].value_counts()
```

Out[45]:

```
-1    4980
 3      5
 2      5
 1      5
 0      5
```

Name: clusters, dtype: int64

In [46]:

```
dbscan = DBSCAN(eps = 0.05)
dbscan.fit(avg_w2v_std_train_data)
final_data['clusters'] = dbscan.labels_
final_data['clusters'].value_counts()
```

Out[46]:

```
-1    4995
 0      5
```

Name: clusters, dtype: int64

In [47]:

```
dbscan = DBSCAN(eps = 10)
dbscan.fit(avg_w2v_std_train_data)
final_data['clusters'] = dbscan.labels_
final_data['clusters'].value_counts()
```

Out[47]:

```
-1    3344
 0    1641
 3       6
 1       6
 2       3
```

Name: clusters, dtype: int64

In [48]:

```
dbscan = DBSCAN(eps = 25)
dbscan.fit(avg_w2v_std_train_data)
final_data['clusters'] = dbscan.labels_
final_data['clusters'].value_counts()
```

Out[48]:

```
 0    4985
-1     15
```

Name: clusters, dtype: int64

5.2) TF-IDF W2V

In [4]:

```
import pickle
with open("tfidf_w2v_train_data.pkl", "rb") as f:
    tfidf_w2v_train_data = pickle.load(f)

with open("tfidf_w2v_test_data.pkl", "rb") as f:
    tfidf_w2v_test_data = pickle.load(f)
```

```
In [5]:
```

```
tfidf_w2v_train_data1 = tfidf_w2v_train_data[0:5000]
```

```
In [6]:
```

```
standardizing = StandardScaler(with_mean = False)
tfidf_w2v_std_train_data = standardizing.fit_transform(tfidf_w2v_train_data1)
tfidf_w2v_std_train_data.shape
```

```
Out[6]:
```

```
(5000, 300)
```

```
In [9]:
```

```
from sklearn.cluster import DBSCAN
dbscan = DBSCAN(eps = 0.005)
dbscan.fit(tfidf_w2v_std_train_data)
final_data['clusters'] = dbscan.labels_
final_data['clusters'].value_counts()
```

```
Out[9]:
```

```
-1    4995
 0         5
```

```
Name: clusters, dtype: int64
```

```
In [10]:
```

```
dbscan = DBSCAN(eps = 0.05)
dbscan.fit(tfidf_w2v_std_train_data)
final_data['clusters'] = dbscan.labels_
final_data['clusters'].value_counts()
```

```
Out[10]:
```

```
-1    4995
 0         5
```

```
Name: clusters, dtype: int64
```

```
In [11]:
```

```
dbscan = DBSCAN(eps = 0.5)
dbscan.fit(tfidf_w2v_std_train_data)
final_data['clusters'] = dbscan.labels_
final_data['clusters'].value_counts()
```

```
Out[11]:
```

```
-1    4995
 0         5
```

```
Name: clusters, dtype: int64
```

```
In [12]:
```

```
dbscan = DBSCAN(eps = 5)
dbscan.fit(tfidf_w2v_std_train_data)
final_data['clusters'] = dbscan.labels_
final_data['clusters'].value_counts()
```

```
Out[12]:
```

```
-1    4903
 0      77
 3       5
 2       5
 1       5
 4       5
```

Name: clusters, dtype: int64

In [13]:

```
dbscan = DBSCAN(eps = 50)
dbscan.fit(tfidf_w2v_std_train_data)
final_data['clusters'] = dbscan.labels_
final_data['clusters'].value_counts()
```

Out[13]:

```
0      5000
Name: clusters, dtype: int64
```

In [15]:

```
dbscan = DBSCAN(eps =0.5,min_samples=5)
dbscan.fit(tfidf_w2v_std_train_data)
final_data['clusters'] = dbscan.labels_
final_data['clusters'].value_counts()
```

Out[15]:

```
-1      4995
0         5
Name: clusters, dtype: int64
```

In [16]:

```
dbscan = DBSCAN(eps =15,min_samples=5)
dbscan.fit(tfidf_w2v_std_train_data)
final_data['clusters'] = dbscan.labels_
final_data['clusters'].value_counts()
```

Out[16]:

```
0      4684
-1       316
Name: clusters, dtype: int64
```

In [18]:

```
dbscan = DBSCAN(eps =5,min_samples=5)
dbscan.fit(tfidf_w2v_std_train_data)
final_data['clusters'] = dbscan.labels_
final_data['clusters'].value_counts()
```

Out[18]:

```
-1      4903
0         77
3         5
2         5
1         5
4         5
Name: clusters, dtype: int64
```

In [17]:

```
dbscan = DBSCAN(eps =100,min_samples=10)
dbscan.fit(tfidf_w2v_std_train_data)
final_data['clusters'] = dbscan.labels_
final_data['clusters'].value_counts()
```

Out[17]:

```
0      5000
Name: clusters, dtype: int64
```


In [20]:

Out[20]:

Name: clusters, dtype: int64

In [21]:

*****word cloud for cluster 0 *****



zhejiang provinc still proud leav wither date fav



*****word cloud for cluster 2 *****



*****word cloud for cluster 3*****

