

# AMAZON LOGISTIC REGRESSION

January 27, 2019

## 1 Amazon Fine Food Reviews Analysis

Data Source: <https://www.kaggle.com/snap/amazon-fine-food-reviews>

EDA: <https://nycdatascience.com/blog/student-works/amazon-fine-foods-visualization/>

The Amazon Fine Food Reviews dataset consists of reviews of fine foods from Amazon.

Number of reviews: 568,454 Number of users: 256,059 Number of products: 74,258 Timespan: Oct 1999 - Oct 2012 Number of Attributes/Columns in data: 10

Attribute Information:

1. Id
2. ProductId - unique identifier for the product
3. UserId - unique identifier for the user
4. ProfileName
5. HelpfulnessNumerator - number of users who found the review helpful
6. HelpfulnessDenominator - number of users who indicated whether they found the review helpful or not
7. Score - rating between 1 and 5
8. Time - timestamp for the review
9. Summary - brief summary of the review
10. Text - text of the review

**Objective:** Given a review, determine whether the review is positive (rating of 4 or 5) or negative (rating of 1 or 2).

[Q] How to determine if a review is positive or negative? [Ans] We could use Score/Rating. A rating of 4 or 5 can be considered as a positive review. A rating of 1 or 2 can be considered as negative one. A review of rating 3 is considered neutral and such reviews are ignored from our analysis. This is an approximate and proxy way of determining the polarity (positivity/negativity) of a review.

## 2 [1]. Reading Data

### 3 [1.1] Loading the data

The dataset is available in two forms 1. .csv file 2. SQLite Database

In order to load the data, We have used the SQLite dataset as it is easier to query the data and visualise the data efficiently.

Here as we only want to get the global sentiment of the recommendations (positive or negative), we will purposefully ignore all Scores equal to 3. If the score is above 3, then the recommendation will be set to "positive". Otherwise, it will be set to "negative".

```
In [1]: import sqlite3
import pandas as pd
import numpy as np
import nltk
import string
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.feature_extraction.text import TfidfTransformer
from sklearn.feature_extraction.text import TfidfVectorizer

from sklearn.feature_extraction.text import CountVectorizer
from sklearn.metrics import confusion_matrix
from sklearn import metrics
from sklearn.metrics import roc_curve, auc

# Tutorial about Python regular expressions: https://pymotw.com/2/re/
import re
import string
from nltk.corpus import stopwords          #importing stopwords
from nltk.stem import SnowballStemmer
from nltk.stem.wordnet import WordNetLemmatizer

from sklearn.model_selection import train_test_split
from collections import Counter
from sklearn.metrics import accuracy_score
from sklearn.linear_model import LogisticRegression
from sklearn.preprocessing import StandardScaler
from sklearn.model_selection import GridSearchCV

from sklearn.metrics import confusion_matrix
from sklearn.metrics import f1_score
from sklearn.metrics import precision_score
from sklearn.metrics import recall_score
import math
import pickle
```

```
In [2]: data = pd.read_csv('Reviews.csv')
print (data.head(2))
print(data.shape)
```

	Id	ProductId	UserId	ProfileName	HelpfulnessNumerator	\
0	1	B001E4KFG0	A3SGXH7AUHU8GW	delmartian	1	
1	2	B00813GRG4	A1D87F6ZCVE5NK	dll pa	0	

	HelpfulnessDenominator	Score	Time	Summary \
0	1	5	1303862400	Good Quality Dog Food
1	0	1	1346976000	Not as Advertised

	Text
0	I have bought several of the Vitality canned d...
1	Product arrived labeled as Jumbo Salted Peanut...

(568454, 10)

### 3.1 [2] Data Cleaning: Deduplication and Nan features

It is observed (as shown in the table below) that the reviews data had many duplicate entries. Hence it was necessary to remove duplicates in order to get unbiased results for the analysis of the data. Following is an example:

```
In [3]: #checking for Nan values in data. True indicates Nan values are present along the column
data.isnull().any()
```

```
Out[3]: Id                False
        ProductId         False
        UserId            False
        ProfileName        True
        HelpfulnessNumerator False
        HelpfulnessDenominator False
        Score              False
        Time               False
        Summary            True
        Text               False
        dtype: bool
```

```
In [4]: # checking for Nan values along 'profilename' column
#data[data['ProfileName'].isnull()].head()
```

```
In [5]: # checking for Nan values along 'summary' column
#data[data['Summary'].isnull()]
```

```
In [6]: #Dropping Nan values
data = data.dropna()
```

```
In [7]: #printing shape of data after dropping Nan values
print (data.shape)
```

(568411, 10)

```
In [8]: #Review score should lie between 1 to 5
#Returns True if all the scores lie between 1 to 5(inclusive)
list1 = data['Score'].map(lambda x: True if x in [1,2,3,4,5] else False)
list1.all()
```

Out[8]: True

```
In [9]: filtered_data = data.loc[data['Score']!=3].copy()
        #print (filtered_data.head())
        print (filtered_data.shape)
```

(525773, 10)

```
In [10]: #mapping positive(>3) and negative(<3) reviews based on scores of the data.
        pos_negative = filtered_data['Score'].map(lambda x: 1 if int (x)>3 else 0)
        filtered_data['Score'] = pos_negative
        print ('shape of filtered_data')
        print (filtered_data.shape)
        print (filtered_data.head(2))
```

shape of filtered\_data

(525773, 10)

	Id	ProductId	UserId	ProfileName	HelpfulnessNumerator	\
0	1	B001E4KFG0	A3SGXH7AUHU8GW	delmartian	1	
1	2	B00813GRG4	A1D87F6ZCVE5NK	dll pa	0	

	HelpfulnessDenominator	Score	Time	Summary	\
0	1	1	1303862400	Good Quality Dog Food	
1	0	0	1346976000	Not as Advertised	

Text

0	I have bought several of the Vitality canned d...
1	Product arrived labeled as Jumbo Salted Peanut...

```
In [11]: #arranging data with increasing productid
        sorted_data = filtered_data.sort_values('ProductId',axis=0,ascending=True,inplace=False)
```

```
In [12]: #finding the duplicates in our data
        #If the same person gives for the same product at the same time we call it as suplica
        #sorted_data.loc[sorted_data.duplicated(["UserId","ProfileName","Time","Text"],keep =
```

```
In [13]: #counting number of duplicates present in our data
        sorted_data.duplicated(["UserId","ProfileName","Time","Text"]).sum()
```

Out[13]: 161612

```
In [14]: #dropping all duplicates keeping the first one
        final = sorted_data.drop_duplicates(subset={"UserId","ProfileName","Time","Text"},keep =
        final.shape
```

Out[14]: (364161, 10)

```

In [15]: #helpfulness numerator denotes number of people who found the review helpful
         #helpfulness denominator denotes number of people who indicated whether or not the re
         #so, helpfulness numerator should be less than denominator
         final = final[final.HelpfulnessNumerator<=final.HelpfulnessDenominator]

In [16]: #final shape of data after preprocessing
         final.shape

Out[16]: (364159, 10)

In [17]: #arranging data with increasing productid
         final = final.sort_values('Time',axis=0,ascending=True,inplace=False,kind='quicksort',

In [18]: final['Score'].value_counts()

Out[18]: 1    307054
         0     57105
         Name: Score, dtype: int64

In [19]: final_data = final.iloc[0:100000,:].copy()

In [20]: final_data.shape

Out[20]: (100000, 10)

In [21]: final_data['Score'].value_counts()

Out[21]: 1     87730
         0    12270
         Name: Score, dtype: int64

```

## 4 [3] Preprocessing

### 5 [3.1]. Preprocessing Review Text

Now that we have finished deduplication our data requires some preprocessing before we go on further with analysis and making the prediction model.

Hence in the Preprocessing phase we do the following in the order below:-

1. Begin by removing the html tags
2. Remove any punctuations or limited set of special characters like , or . or # etc.
3. Check if the word is made up of english letters and is not alpha-numeric
4. Check to see if the length of the word is greater than 2 (as it was researched that there is no adjective in 2-letters)
5. Convert the word to lowercase
6. Remove Stopwords
7. Finally Snowball Stemming the word (it was observed to be better than Porter Stemming)

After which we collect the words used to describe positive and negative reviews

```

In [98]: stop = set(stopwords.words('english')) #set of stopwords
sno = nltk.stem.SnowballStemmer('english') #initialising the snowball stemmer
not_words = re.findall(r'\w*n[\'|o]t',str (stop)) #finding NOT words in stop words
not_words.append('n\'t')
not_words.append('no')
print (not_words)
stop_words = stop - set (not_words) #removing NOT words from stop words
# https://stackoverflow.com/a/47091490/4084039
import re

def decontracted(phrase):
    # specific
    phrase = re.sub(r'\w*n[\'|o]t', "not", phrase)
    # general
    phrase = re.sub(r"\ 're", " are", phrase)
    phrase = re.sub(r"\ 's", " is", phrase)
    phrase = re.sub(r"\ 'd", " would", phrase)
    phrase = re.sub(r"\ 'll", " will", phrase)
    phrase = re.sub(r"\ 't", " not", phrase)
    phrase = re.sub(r"\ 've", " have", phrase)
    phrase = re.sub(r"\ 'm", " am", phrase)
    return phrase

def cleanhtmlpunc(sentence): #function to clean the word of any html-tags
    clean = re.compile('<.*?>')
    clean = re.sub(clean, ' ', sentence)
    clean = re.sub(r"(http|www)\S+", "", clean)
    clean = re.sub(r"\S+com", "",clean)
    #clean = re.sub(r"(\w+)", "",clean)
    clean = re.sub(r"\.", " ",clean)
    cleaned = re.sub(r'[?+!+|\'+|"+|#+|:|+]',r' ',clean)
    cleantext = re.sub(r'[\.,+|+|+|+|+|/+|/]',r' ',cleaned)
    return cleantext

#def cleanpunc(sentence): #function to clean the word of any punctuation or special c
    # return cleaned
print('*****')
print(stop_words)

["weren't", "didn't", "won't", "aren't", "hadn't", "wouldn't", "mustn't", "couldn't", "mightn't",
*****
{'how', 'before', 'yourselves', 'if', 'is', 'have', 'during', 'same', 'our', 'some', 'hers', 'v

```

```

In [23]: def cleanedtext(reviews):
    str1=' '
    final_string=[]
    s=' '
    for sent in reviews:
        filtered_sentence=[]

```

```

sent=cleanhtmlpunc(decontracted(sent)) # remove HTML tags
for w in sent.split():
    for cleaned_words in w.split():
        if((cleaned_words.isalpha()) & (len(cleaned_words)>2)):
            if((cleaned_words.lower() not in stop_words)):
                s=(sno.stem(cleaned_words.lower())).encode('utf8')
                filtered_sentence.append(s)
            else:
                continue
        else:
            continue
    str1 = b" ".join(filtered_sentence) #final string of cleaned words for review
    final_string.append(str1)
return final_string

```

```
In [24]: final_string = cleanedtext(final_data['Text'].values)
```

## 6 Applying Logistic Regression

### 7 [4.1] Logistic Regression on BOW, SET 1

#### 7.1 [4.1.1] Applying Logistic Regression with L2 regularization on BOW, SET 1

```

In [2]: import pickle
        with open("final_data.pkl", "rb") as f:
            final_data = pickle.load(f)

In [3]: final_data = final_data.iloc[0:100000,:]
```

In [4]: final\_data.shape

Out[4]: (100000, 13)

```

In [5]: #https://stackoverflow.com/questions/38640109/logistic-regression-python-solvers-defin

In [5]: # split the data set into train and test
        X_train, X_test, y_train, y_test = train_test_split(final_data['CleanedText'].values, ,
                                                            test_size=0.3, random_sta

In [6]: model = CountVectorizer(dtype=float)
        final_counts= model.fit_transform(X_train)

In [7]: #standardizing the bag of words
        standardizing = StandardScaler(with_mean = False)
        final_std_data = standardizing.fit_transform(final_counts)
        final_std_data.shape

Out[7]: (70000, 31094)

```

```

In [10]: from sklearn.metrics import make_scorer
         from sklearn.model_selection import cross_val_score
         from sklearn.metrics import roc_auc_score
         # creating list for hyperparameter alpha
         alpha_values = [0.0001,0.0005,0.001,0.005,0.01,0.05,0.1,100,500,1000,2500,5000,10000]
         # empty list that will hold cv scores

         cv_scores = []
         train_auc_values = []

         # perform 10-fold cross validation
         for alpha in alpha_values:
             log_reg = LogisticRegression(C=alpha,solver='saga',max_iter=10000)
             auc = make_scorer(roc_auc_score,greater_is_better=True,
                              needs_threshold=True)
             auc_scores = cross_val_score(log_reg, final_std_data, y_train, cv=10, scoring=auc)
             log_reg.fit(final_std_data,y_train)
             y_pred_proba = log_reg.predict_proba(final_std_data)[:,:1]
             train_auc = roc_auc_score(y_train, y_pred_proba)
             train_auc_values.append(train_auc)
             cv_scores.append(auc_scores.mean())

         print ('train scores')
         print (train_auc_values)
         print (' '*50)
         print ('CV scores')
         print (cv_scores)


         # changing to misclassification error

         log = [math.log10(x) for x in alpha_values]
         # plot misclassification error vs alpha
         plt.plot(log, cv_scores,label='CV')
         #plt.label('cv_f1')
         plt.plot(log,train_auc_values,label='train')
         #plt.label('train_f1')
         plt.legend()
         plt.xlabel('hyperparameter C')
         plt.ylabel('AUC')
         plt.show()

train scores
[0.8872166258768911, 0.9032958584075139, 0.9055026416848345, 0.9060861740180313, 0.90618246203
*****
CV scores
[0.7675376832793198, 0.786533751980348, 0.7891311431371723, 0.791253547700876, 0.7912544583195

```





```
In [11]: log_reg = LogisticRegression(C=100,solver='saga',max_iter=10000,tol=0.001,n_jobs=-1)
log_reg.fit(final_counts,y_train)
```

```
Out[11]: LogisticRegression(C=100, class_weight=None, dual=False, fit_intercept=True,
intercept_scaling=1, max_iter=10000, multi_class='ovr',
n_jobs=-1, penalty='l2', random_state=None, solver='saga',
tol=0.001, verbose=0, warm_start=False)
```

```
In [12]: print('accuracy = {0}'.format(log_reg.score(standardizing.transform(model.transform(X_test)),
accuracy = 92.5
```


```
In [13]: predictions = log_reg.predict(standardizing.transform(model.transform(X_test)))
```

```
In [14]: precision = precision_score(y_test, predictions,pos_label=1)
Recall = recall_score(y_test, predictions,pos_label=1)
f1 = f1_score(y_test, predictions,pos_label=1)
print ('precision = {0}'.format(precision))
print ('Recall={0}'.format(Recall))
print ('f1_score={0}'.format(f1))
```


```
precision = 0.9666765843498959
Recall=0.8781081081081081
f1_score=0.9202662512392012
```

```
In [15]: from sklearn.metrics import confusion_matrix
import seaborn as sns
result = confusion_matrix(y_test,predictions)
#print(result)
sns.set(font_scale=1.4)#for label size
sns.heatmap(result, annot=True,annot_kws={"size": 16}, fmt='g')
```

Out[15]: <matplotlib.axes.\_subplots.AxesSubplot at 0x4496e6f6d8>



```
In [16]: y_pred_proba = log_reg.predict_proba(model.transform(X_test))[:,1]
fpr, tpr, _ = roc_curve(y_test, y_pred_proba,pos_label=1 )
#auc = roc_auc_score(y_test, y_pred_proba)
auc = np.trapz(tpr,fpr)
plt.plot(fpr,tpr,label="data 1, auc="+str(auc))
plt.legend(loc=4)
plt.show()
```



### 7.1.1 [4.1.1.1] Pertubation Test

```
In [17]: #clf = LogisticRegression(C=model.best_params_['C'], penalty='l2')
#clf.fit(final_std_data, y_train)
weights_before = log_reg.coef_
print ('weights_before={0}'.format(weights_before))
```

```
weights_before=[[ 1.67703901e-02  5.14711165e-03 -5.38907139e-02 ...  6.79927908e-03
 6.29820289e-05  6.85599065e-04]]
```

```
In [18]: from scipy.sparse import find
eps = np.random.normal(scale=0.01)
print ('noise adding to weight vector is = {0}'.format(eps))
X_train1 = final_counts
a,b,c = find(X_train1)
X_train1[a,b] = X_train1[a,b]+eps
log_reg.fit(X_train1,y_train)
weights_after = log_reg.coef_
print ('weights_after adding noise = {0}'.format(weights_after))
```

```
noise adding to weight vector is = 0.006139962498737992
```

```
weights_after adding noise = [[ 1.66861339e-02  5.18499791e-03 -5.34311971e-02 ...  6.90404251e-03
 4.97940225e-05  7.01490370e-04]]
```

```

In [19]: per_weights_diff = abs((weights_before - weights_after)/weights_before)*100
         per_weights_diff = per_weights_diff[0]

In [20]: for i in range(0,110,10):
         print ('{0} percentile value is {1}'.format(i,np.percentile(per_weights_diff,i)))

0 percentile value is 0.00016873855680106447
10 percentile value is 0.17566958318672576
20 percentile value is 0.3708582422088057
30 percentile value is 0.588400474177819
40 percentile value is 0.8893549312443004
50 percentile value is 1.3493502065878444
60 percentile value is 2.144811872848045
70 percentile value is 3.7624360694869883
80 percentile value is 7.594158041619698
90 percentile value is 21.302181988035613
100 percentile value is 1359298.8626627645

In [21]: for i in range(0,11,1):
         print ('{0} percentile value is {1}'.format(60+i,np.percentile(per_weights_diff,i)))

60 percentile value is 2.144811872848045
61 percentile value is 2.254222670672485
62 percentile value is 2.370316193944895
63 percentile value is 2.494914457111536
64 percentile value is 2.6422856736299014
65 percentile value is 2.786969047610154
66 percentile value is 2.959162340942017
67 percentile value is 3.1436518761189367
68 percentile value is 3.3479251476185516
69 percentile value is 3.5527140753665396
70 percentile value is 3.7624360694869883

In [22]: for i in range(0,11,1):
         print ('{0} percentile value is {1}'.format(63+i*0.1,np.percentile(per_weights_d

63.0 percentile value is 2.494914457111536
63.1 percentile value is 2.513683703079929
63.2 percentile value is 2.53208235928024
63.3 percentile value is 2.5436407867107893
63.4 percentile value is 2.5554346956896175
63.5 percentile value is 2.573439426220451
63.6 percentile value is 2.5871164719424375
63.7 percentile value is 2.6020419986022927
63.8 percentile value is 2.6134030229436798
63.9 percentile value is 2.629241177161071
64.0 percentile value is 2.6422856736299014

```

```
In [23]: feat_names = model.get_feature_names()

In [24]: per_weights_diff[per_weights_diff>2.5].size

Out[24]: 11492

In [25]: (per_weights_diff[per_weights_diff>2.5].size)*100/weights_before.size

Out[25]: 36.958898822924034

In [26]: index = per_weights_diff>2.5

In [27]: colli_features = [feature for feature, bol in zip(feat_names, index) if bol==True]

In [28]: colli_features[0:10]

Out[28]: ['aaaah',
          'aaah',
          'aad',
          'aadp',
          'aahh',
          'ab',
          'abandon',
          'abaolut',
          'abba',
          'abbi']
```

### 7.1.2 [4.1.1.2] Feature Importance

```
In [29]: indices = list(np.argsort(weights_before[0]))

In [30]: pos_indices = indices[-10:]
neg_indices = indices[0:10]

In [31]: print('top 10 positive features\t\t\ttop 10 negative features')
print('-'*100)
for id1,id2 in zip(pos_indices,neg_indices):
    print('\t{0}\t\t\t\t\t\t\t{1}'.format(feats_names[id1],feats_names[id2]))
```

top 10 positive features

top 10 negative features

awesom  
fantast  
yummi  
addict  
great  
amaz  
best  
perfect  
excel  
delici

worst  
terribl  
aw  
bland  
horribl  
disappoint  
threw  
unfortun  
stale  
tasteless

## 7.2 [4.1.2] Applying Logistic Regression with L1 regularization on BOW, SET 1

```
In [32]: from sklearn.metrics import make_scorer
         from sklearn.model_selection import cross_val_score
         from sklearn.metrics import roc_auc_score
         # creating list for hyperparameter alpha
         alpha_values = [0.0001,0.0005,0.001,0.005,0.01,0.05,0.1,100,500,1000,2500,5000,10000]
         # empty list that will hold cv scores

         cv_scores = []
         train_auc_values = []


         # perform 10-fold cross validation
         for alpha in alpha_values:
             log_reg = LogisticRegression(C=alpha,penalty='l1',max_iter=10000)
             auc = make_scorer(roc_auc_score,greater_is_better=True,
                               needs_threshold=True)
             auc_scores = cross_val_score(log_reg, final_std_data, y_train, cv=10, scoring=auc)
             log_reg.fit(final_std_data,y_train)
             y_pred_proba = log_reg.predict_proba(final_std_data)[:,:1]
             train_auc = roc_auc_score(y_train, y_pred_proba)
             train_auc_values.append(train_auc)
             cv_scores.append(auc_scores.mean())

         print ('train scores')
         print (train_auc_values)
         print (''*50)
         print ('CV scores')
         print (cv_scores)

         # changing to misclassification error

         log = [math.log10(x) for x in alpha_values]
         # plot misclassification error vs alpha
         plt.plot(log, cv_scores,label='CV')
         #plt.label('cv_f1')
         plt.plot(log,train_auc_values,label='train')
         #plt.label('train_f1')
         plt.legend()
         plt.xlabel('hyperparameter C')
         plt.ylabel('AUC')
         plt.show()

train scores
[0.5, 0.7940562566329163, 0.8592942388920859, 0.9422173038918453, 0.9645540175998825, 0.988524
*****
CV scores
[0.5, 0.719741226931902, 0.8386816934576604, 0.9158721469675113, 0.9254130744246624, 0.9205910
```



In [ ]: <http://rnowling.github.io/data/science/2016/09/04/comparing-lr-regularization-and-opt>

```
In [33]: log_reg = LogisticRegression(C=0.005,penalty='l1',max_iter=5000,tol=0.001)
log_reg.fit(final_counts,y_train)
```

```
Out[33]: LogisticRegression(C=0.005, class_weight=None, dual=False, fit_intercept=True,
intercept_scaling=1, max_iter=5000, multi_class='ovr', n_jobs=1,
penalty='l1', random_state=None, solver='liblinear', tol=0.001,
verbose=0, warm_start=False)
```

```
In [34]: predictions = log_reg.predict(standardizing.transform(model.transform(X_test)))
```


```
In [36]: print('accuracy = {0}'.format(log_reg.score(standardizing.transform(model.transform(X_
accuracy = 87.34666666666666
```

```
In [37]: precision = precision_score(y_test, predictions,pos_label=1)
Recall = recall_score(y_test, predictions,pos_label=1)
f1 = f1_score(y_test, predictions,pos_label=1)
print ('precision = {0}'.format(precision))
print ('Recall={0}'.format(Recall))
print ('f1_score={0}'.format(f1))
```

```
precision = 0.8935617121287658
Recall=0.9688416988416988
f1_score=0.9296802637916342
```


```
In [38]: from sklearn.metrics import confusion_matrix
import seaborn as sns
result = confusion_matrix(y_test,predictions)
#print(result)
sns.set(font_scale=1.4)#for label size
sns.heatmap(result, annot=True,annot_kws={"size": 16}, fmt='g')
```

```
Out[38]: <matplotlib.axes._subplots.AxesSubplot at 0x44818f5940>
```



```
In [39]: y_pred_proba = log_reg.predict_proba(model.transform(X_test))[:,1]
fpr, tpr, _ = roc_curve(y_test, y_pred_proba,pos_label=1 )
#auc = roc_auc_score(y_test, y_pred_proba)
auc = np.trapz(tpr,fpr)
plt.plot(fpr,tpr,label="data 1, auc="+str(auc))
plt.legend(loc=4)
plt.show()
```





### 7.2.1 [4.1.2.1] Pertubation Test

```
In [40]: from scipy.sparse import find
weights_before = log_reg.coef_
print ('weights_before={0}'.format(weights_before))
eps = np.random.normal(scale=0.01)
print ('noise adding to weight vector is = {0}'.format(eps))
X_train1 = final_counts
a,b,c = find(X_train1)
X_train1[a,b] = X_train1[a,b]+eps
log_reg.fit(X_train1,y_train)
weights_after = log_reg.coef_
print ('weights_after adding noise = {0}'.format(weights_after))

weights_before=[[0. 0. 0. ... 0. 0. 0.]]
noise adding to weight vector is = -0.013766085204073575
weights_after adding noise = [[0. 0. 0. ... 0. 0. 0.]]

In [41]: weights_before = weights_before[0]+0.000001

In [42]: weights_after = weights_after[0]+0.000001

In [43]: per_weights_diff = abs((weights_before - weights_after)/weights_before)*100
per_weights_diff = per_weights_diff[0]
```

```
In [44]: for i in range(0,110,10):
          print ('{0} percentile value is {1}'.format(i,np.percentile(per_weights_diff,i))

0 percentile value is 0.0
10 percentile value is 0.0
20 percentile value is 0.0
30 percentile value is 0.0
40 percentile value is 0.0
50 percentile value is 0.0
60 percentile value is 0.0
70 percentile value is 0.0
80 percentile value is 0.0
90 percentile value is 0.0
100 percentile value is 0.0
```

```
In [45]: (per_weights_diff[per_weights_diff>2.5].size)*100/weights_before.size
Out[45]: 0.0
```

## 7.2.2 [4.1.2.2] Calculating sparsity on L1 regularization

```
In [46]: print ('Hyper Parameters(C) \t Num of non Zero elements')
          alpha_values = [0.0001,0.0005,0.001,0.005,0.01,0.05,0.1,100,500,1000,2500,5000,10000]
          for c in alpha_values:
              clf = LogisticRegression(C=c, penalty='l1',tol=0.01)
              clf.fit(final_counts, y_train)
              w = clf.coef_
              print ('\t{0}\t\t{1}'.format(c,np.count_nonzero(w)))
```

Hyper Parameters(C)	Num of non Zero elements
0.0001	0
0.0005	2
0.001	7
0.005	51
0.01	118
0.05	451
0.1	744
100	13555
500	15953
1000	16784
2500	19333
5000	20597
10000	23894

## 7.2.3 [4.1.2.3] Feature Importance

```
In [48]: feat_names = model.get_feature_names()
          weights_before = log_reg.coef_
```

top 10 positive features                      top 10 negative features

## 8 [4.2] Logistic Regression on TFIDF, SET 2

```

train_auc_values = []

# perform 10-fold cross validation
for alpha in alpha_values:
    log_reg = LogisticRegression(C=alpha,penalty='l2',max_iter=10000)
    auc = make_scorer(roc_auc_score,greater_is_better=True,
                      needs_threshold=True)
    auc_scores = cross_val_score(log_reg, final_std_data, y_train, cv=10, scoring=auc)
    log_reg.fit(final_std_data,y_train)
    y_pred_proba = log_reg.predict_proba(final_std_data)[:,:1]
    train_auc = roc_auc_score(y_train, y_pred_proba)
    train_auc_values.append(train_auc)
    cv_scores.append(auc_scores.mean())

print ('train scores')
print (train_auc_values)
print ('*'*50)
print ('CV scores')
print (cv_scores)

# changing to misclassification error

log = [math.log10(x) for x in alpha_values]
# plot misclassification error vs alpha
plt.plot(log, cv_scores,label='CV')
#plt.label('cv_f1')
plt.plot(log,train_auc_values,label='train')
#plt.label('train_f1')
plt.legend()
plt.xlabel('hyperparameter C')
plt.ylabel('AUC')
plt.show()

```


train scores

[0.9801327058379166, 0.989632704947094, 0.9918192724909438, 0.9947721691588912, 0.995414761048

\*\*\*\*\*

CV scores

[0.9092294068052114, 0.9140867851223128, 0.910187268720191, 0.893801488307162, 0.8852491462455



```
In [62]: log_reg = LogisticRegression(C=0.0005,max_iter=10000,tol=0.001)
log_reg.fit(final_std_data,y_train)
```

```
Out[62]: LogisticRegression(C=0.0005, class_weight=None, dual=False,
fit_intercept=True, intercept_scaling=1, max_iter=10000,
multi_class='ovr', n_jobs=1, penalty='l2', random_state=None,
solver='liblinear', tol=0.001, verbose=0, warm_start=False)
```

```
In [63]: print('accuracy = {}'.format(log_reg.score(standardizing.transform(vector.transform(
accuracy = 90.47
```


```
In [64]: predictions = log_reg.predict(standardizing.transform(vector.transform(tfidf_test)))
```

```
In [65]: precision = precision_score(y_test, predictions,pos_label=1)
Recall = recall_score(y_test, predictions,pos_label=1)
f1 = f1_score(y_test, predictions,pos_label=1)
print ('precision = {}'.format(precision*100))
print ('Recall={}'.format(Recall*100))
print ('f1_score={}'.format(f1*100))
```


```
precision = 92.23831347387717
Recall=97.13513513513513
f1_score=94.62341325811
```

```
In [66]: from sklearn.metrics import confusion_matrix
import seaborn as sns
result = confusion_matrix(y_test,predictions)
#print(result)
sns.set(font_scale=1.4)#for label size
sns.heatmap(result, annot=True,annot_kws={"size": 16}, fmt='g')
```

Out[66]: <matplotlib.axes.\_subplots.AxesSubplot at 0x4485e46da0>



```
In [67]: y_pred_proba = log_reg.predict_proba(standardizing.transform(vector.transform(tfidf_t
fpr, tpr, _ = roc_curve(y_test, y_pred_proba,pos_label=1 )
#auc = roc_auc_score(y_test, y_pred_proba)
auc = np.trapz(tpr,fpr)
plt.plot(fpr,tpr,label="data 1, auc="+str(auc))
plt.legend(loc=4)
plt.show()
```



### 8.0.1 [4.2.1.1] Feature Importance of TFIDF with l2 Regularizer

```
In [68]: pos_indices = []
neg_indices = []
feat_names = vector.get_feature_names()
weights_before = log_reg.coef_
indices = list(np.argsort(weights_before[0]))
pos_indices = indices[-10:]
neg_indices = indices[0:10]
print('top 10 positive features\t\t\ttop 10 negative features')
print('-'*100)
for id1,id2 in zip(pos_indices,neg_indices):
    print('\t{0}\t\t\t\t\t{1}'.format(feat_names[id1],feat_names[id2]))
```

top 10 positive features

top 10 negative features

find  
nice  
favorit  
excel  
perfect  
delici  
good  
best

not  
disappoint  
worst  
aw  
terribl  
horribl  
unfortun  
stale

love	bland
great	return

## 8.0.2 [4.2.2] Applying Logistic Regression with L1 regularization on TFIDE, SET 2

```
In [69]: from sklearn.metrics import make_scorer
from sklearn.model_selection import cross_val_score
from sklearn.metrics import roc_auc_score
# creating list for hyperparameter alpha
alpha_values = [0.0001,0.0005,0.001,0.005,0.01,0.05,0.1,100,500,1000,2500,5000,10000]
# empty list that will hold cv scores

cv_scores = []
train_auc_values = []

# perform 10-fold cross validation
for alpha in alpha_values:
    log_reg = LogisticRegression(C=alpha,penalty='l1',max_iter=10000)
    auc = make_scorer(roc_auc_score,greater_is_better=True,
                      needs_threshold=True)
    auc_scores = cross_val_score(log_reg, final_std_data, y_train, cv=10, scoring=auc)
    log_reg.fit(final_std_data,y_train)
    y_pred_proba = log_reg.predict_proba(final_std_data)[:,:1]
    train_auc = roc_auc_score(y_train, y_pred_proba)
    train_auc_values.append(train_auc)
    cv_scores.append(auc_scores.mean())

print ('train scores')
print (train_auc_values)
print ('*'*50)
print ('CV scores')
print (cv_scores)

# changing to misclassification error

log = [math.log10(x) for x in alpha_values]
# plot misclassification error vs alpha
plt.plot(log, cv_scores,label='CV')
#plt.label('cv_f1')
plt.plot(log,train_auc_values,label='train')
#plt.label('train_f1')
plt.legend()
plt.xlabel('hyperparameter C')
plt.ylabel('AUC')
plt.show()
```

train scores




[0.5, 0.7909062733902787, 0.8726965565352625, 0.9483963491319725, 0.9687353694765786, 0.989875

\*\*\*\*\*

CV scores

[0.5, 0.737474643725412, 0.8531336020054198, 0.9265739647008588, 0.9360879546733643, 0.9275145



```
In [70]: log_reg = LogisticRegression(C=0.005,penalty='l1',max_iter=10000,tol=0.001)
log_reg.fit(final_std_data,y_train)
```

```
Out[70]: LogisticRegression(C=0.005, class_weight=None, dual=False, fit_intercept=True,
intercept_scaling=1, max_iter=10000, multi_class='ovr', n_jobs=1,
penalty='l1', random_state=None, solver='liblinear', tol=0.001,
verbose=0, warm_start=False)
```

```
In [71]: print('accuracy = {0}'.format(log_reg.score(standardizing.transform(vector.transform(
accuracy = 91.13333333333333
```


```
In [72]: predictions = log_reg.predict(standardizing.transform(vector.transform(tfidf_test)))
```

```
In [73]: precision = precision_score(y_test, predictions,pos_label=1)
Recall = recall_score(y_test, predictions,pos_label=1)
f1 = f1_score(y_test, predictions,pos_label=1)
print ('precision = {0}'.format(precision))
print ('Recall={0}'.format(Recall))
print ('f1_score={0}'.format(f1))
```


```
precision = 0.916547175222254
Recall=0.9871814671814672
f1_score=0.9505539445311919
```

```
In [74]: from sklearn.metrics import confusion_matrix
import seaborn as sns
result = confusion_matrix(y_test,predictions)
#print(result)
sns.set(font_scale=1.4)#for label size
sns.heatmap(result, annot=True,annot_kws={"size": 16}, fmt='g')
```

```
Out[74]: <matplotlib.axes._subplots.AxesSubplot at 0x448149f780>
```



```
In [75]: y_pred_proba = log_reg.predict_proba(standardizing.transform(vector.transform(tfidf_t
fpr, tpr, _ = roc_curve(y_test, y_pred_proba,pos_label=1 )
#auc = roc_auc_score(y_test, y_pred_proba)
auc = np.trapz(tpr,fpr)
plt.plot(fpr,tpr,label="data 1, auc="+str(auc))
plt.legend(loc=4)
plt.show()
```



### 8.0.3 [4.2.2.1] Feature Importance of TFIDF with l1 regularization

```
In [76]: pos_indices = []
neg_indices = []
feat_names = vector.get_feature_names()
weights_before = log_reg.coef_
indices = list(np.argsort(weights_before[0]))
pos_indices = indices[-10:]
neg_indices = indices[0:10]
print('top 10 positive features \t\t top 10 negative features')
print('-'*100)
for id1,id2 in zip(pos_indices,neg_indices):
    print('\t{0}\t\t\t\t\t{1}'.format(feat_names[id1],feat_names[id2]))
```

top 10 positive features

top 10 negative features

find  
nice  
favorit  
excel  
good  
perfect  
delici  
love

not  
disappoint  
worst  
terribl  
aw  
horribl  
return  
money

best  
great

threw  
unfortun

## 9 [4.3] Logistic Regression on AVG W2V, SET 3

### 9.1 [4.3.1] Applying Logistic Regression with L2 regularization on AVG W2V SET 3

```
In [77]: import pickle
         with open("avg_w2v_train_data.pkl", "rb") as f:
             avg_w2v_train_data = pickle.load(f)

In [78]: import pickle
         with open("avg_w2v_test_data.pkl", "rb") as f:
             avg_w2v_test_data = pickle.load(f)

In [79]: standardizing = StandardScaler(with_mean = False)
         avg_w2v_std_train_data = standardizing.fit_transform(avg_w2v_train_data)
         avg_w2v_std_train_data.shape

Out[79]: (70000, 300)

In [80]: from sklearn.metrics import make_scorer
         from sklearn.model_selection import cross_val_score
         from sklearn.metrics import roc_auc_score
         # creating list for hyperparameter alpha
         alpha_values = [0.0001,0.0005,0.001,0.005,0.01,0.05,0.1,100,500,1000,2500,5000,10000]
         # empty list that will hold cv scores

         cv_scores = []
         train_auc_values = []

         # perform 10-fold cross validation
         for alpha in alpha_values:
             log_reg = LogisticRegression(C=alpha,penalty='l2',max_iter=10000)
             auc = make_scorer(roc_auc_score,greater_is_better=True,
                              needs_threshold=True)
             auc_scores = cross_val_score(log_reg, avg_w2v_std_train_data, y_train, cv=10, scor
             log_reg.fit(avg_w2v_std_train_data,y_train)
             y_pred_proba = log_reg.predict_proba(avg_w2v_std_train_data)[:,:1]
             train_auc = roc_auc_score(y_train, y_pred_proba)
             train_auc_values.append(train_auc)
             cv_scores.append(auc_scores.mean())

         print ('train scores')
         print (train_auc_values)
         print ('*'*50)
         print ('CV scores')
```

```

print (cv_scores)

# changing to misclassification error

log = [math.log10(x) for x in alpha_values]
# plot misclassification error vs alpha
plt.plot(log, cv_scores, label='CV')
#plt.label('cv_f1')
plt.plot(log, train_auc_values, label='train')
#plt.label('train_f1')
plt.legend()
plt.xlabel('hyperparameter C')
plt.ylabel('AUC')
plt.show()

```

train scores

```

[0.9039903001300007, 0.9109722605770828, 0.912938977070425, 0.9162205654902067, 0.916893399816
*****


```

CV scores

```

[0.9017952549247147, 0.9087825405111462, 0.910729067005892, 0.9134717315274579, 0.914239264251

```



```

In [81]: log_reg = LogisticRegression(C=0.05,penalty='l2',max_iter=10000,tol=0.001)
log_reg.fit(avg_w2v_std_train_data,y_train)

```


```
Out[81]: LogisticRegression(C=0.05, class_weight=None, dual=False, fit_intercept=True,
                             intercept_scaling=1, max_iter=10000, multi_class='ovr', n_jobs=1,
                             penalty='l2', random_state=None, solver='liblinear', tol=0.001,
                             verbose=0, warm_start=False)
```

```
In [82]: predictions = log_reg.predict(standardizing.transform(avg_w2v_test_data))
print('accuracy = {0}'.format(log_reg.score(standardizing.transform(avg_w2v_test_data),
precision = precision_score(y_test, predictions,pos_label=1)
Recall = recall_score(y_test, predictions,pos_label=1)
f1 = f1_score(y_test, predictions,pos_label=1)
print ('precision = {0}'.format(precision))
print ('Recall={0}'.format(Recall))
print ('f1_score={0}'.format(f1))
```


```
accuracy = 81.02333333333334
precision = 0.9090651443378275
Recall=0.8669111969111969
f1_score=0.8874878950176881
```

```
In [83]: from sklearn.metrics import confusion_matrix
import seaborn as sns
result = confusion_matrix(y_test,predictions)
#print(result)
sns.set(font_scale=1.4)#for label size
sns.heatmap(result, annot=True,annot_kws={"size": 16}, fmt='g')
```

```
Out[83]: <matplotlib.axes._subplots.AxesSubplot at 0x44800a20b8>
```



```
In [84]: y_pred_proba = log_reg.predict_proba(standardizing.transform(avg_w2v_test_data))[:,1]
fpr, tpr, _ = roc_curve(y_test, y_pred_proba, pos_label=1 )
#auc = roc_auc_score(y_test, y_pred_proba)
auc = np.trapz(tpr,fpr)
plt.plot(fpr,tpr,label="data 1, auc="+str(auc))
plt.legend(loc=4)
plt.show()
```



## 9.2 [4.3.2] Applying Logistic Regression with L1 regularization on AVG W2V, SET 3

```
In [85]: from sklearn.metrics import make_scorer
from sklearn.model_selection import cross_val_score
from sklearn.metrics import roc_auc_score
# creating list for hyperparameter alpha
alpha_values = [0.0001,0.0005,0.001,0.005,0.01,0.05,0.1,100,500,1000,2500,5000,10000]
# empty list that will hold cv scores

cv_scores = []
train_auc_values = []

# perform 10-fold cross validation
for alpha in alpha_values:
```

```

log_reg = LogisticRegression(C=alpha,penalty='l1',max_iter=10000,)
auc = make_scorer(roc_auc_score,greater_is_better=True,
                  needs_threshold=True)
auc_scores = cross_val_score(log_reg, avg_w2v_std_train_data, y_train, cv=10, scorer=auc)
log_reg.fit(avg_w2v_std_train_data,y_train)
y_pred_proba = log_reg.predict_proba(avg_w2v_std_train_data)[::,1]
train_auc = roc_auc_score(y_train, y_pred_proba)
train_auc_values.append(train_auc)
cv_scores.append(auc_scores.mean())

print ('train scores')
print (train_auc_values)
print ('*'*50)
print ('CV scores')
print (cv_scores)


# changing to misclassification error

log = [math.log10(x) for x in alpha_values]
# plot misclassification error vs alpha
plt.plot(log, cv_scores,label='CV')
#plt.label('cv_f1')
plt.plot(log,train_auc_values,label='train')
#plt.label('train_f1')
plt.legend()
plt.xlabel('hyperparameter C')
plt.ylabel('AUC')
plt.show()

train scores
[0.596049529536806, 0.8629598876454986, 0.891350149489925, 0.9090660408341188, 0.9120050109759]
*****
CV scores
[0.5127810767906871, 0.8380087759880162, 0.8837906321494697, 0.905839213257182, 0.909713212541]

```





```
In [89]: log_reg = LogisticRegression(C=0.05,penalty='l1',max_iter=10000,tol=0.001)
log_reg.fit(avg_w2v_std_train_data,y_train)
```

```
Out[89]: LogisticRegression(C=0.05, class_weight=None, dual=False, fit_intercept=True,
    intercept_scaling=1, max_iter=10000, multi_class='ovr', n_jobs=1,
    penalty='l1', random_state=None, solver='liblinear', tol=0.001,
    verbose=0, warm_start=False)
```

```
In [90]: predictions = log_reg.predict(standardizing.transform(avg_w2v_test_data))
print('accuracy = {0}'.format(log_reg.score(standardizing.transform(avg_w2v_test_data),
    y_test)))
precision = precision_score(y_test, predictions,pos_label=1)
Recall = recall_score(y_test, predictions,pos_label=1)
f1 = f1_score(y_test, predictions,pos_label=1)
print ('precision = {0}'.format(precision))
print ('Recall={0}'.format(Recall))
print ('f1_score={0}'.format(f1))
```

```
accuracy = 86.45666666666668
```

```
precision = 0.8645941162720807
```

```
Recall=0.9996911196911197
```

```
f1_score=0.9272476587820296
```


```
In [92]: from sklearn.metrics import confusion_matrix
import seaborn as sns
```

```

result = confusion_matrix(y_test,predictions)
#print(result)
sns.set(font_scale=1.4)#for label size
sns.heatmap(result, annot=True,annot_kws={"size": 16}, fmt='g')

```


Out[92]: <matplotlib.axes.\_subplots.AxesSubplot at 0x4485e20358>



```

In [93]: y_pred_proba = log_reg.predict_proba(standardizing.transform(avg_w2v_test_data))[:,1]
fpr, tpr, _ = roc_curve(y_test, y_pred_proba,pos_label=1 )
#auc = roc_auc_score(y_test, y_pred_proba)
auc = np.trapz(tpr,fpr)
plt.plot(fpr,tpr,label="data 1, auc="+str(auc))
plt.legend(loc=4)
plt.show()

```



## 10 [4.4] Logistic Regression on TFIDF W2V, SET 4

### 10.1 [4.4.1] Applying Logistic Regression with L2 regularization on TFIDF W2V, SET 4

```
In [95]: import pickle
         with open("tfidf_w2v_train_data.pkl", "rb") as f:
             tfidf_w2v_train_data = pickle.load(f)

In [96]: import pickle
         with open("tfidf_w2v_test_data.pkl", "rb") as f:
             tfidf_w2v_test_data = pickle.load(f)

In [97]: standardizing = StandardScaler(with_mean = False)
         tfidf_w2v_std_train_data = standardizing.fit_transform(tfidf_w2v_train_data)
         tfidf_w2v_std_train_data.shape

Out[97]: (70000, 300)

In [98]: tfidf_w2v_std_train_data.mean()

Out[98]: 0.0071627058222749704

In [99]: from sklearn.metrics import make_scorer
         from sklearn.model_selection import cross_val_score
```

```

from sklearn.metrics import roc_auc_score
# creating list for hyperparameter alpha
alpha_values = [0.0001,0.0005,0.001,0.005,0.01,0.05,0.1,100,500,1000,2500,5000,10000]
# empty list that will hold cv scores

cv_scores = []
train_auc_values = []

# perform 10-fold cross validation
for alpha in alpha_values:
    log_reg = LogisticRegression(C=alpha,penalty='l2',max_iter=10000)
    auc = make_scorer(roc_auc_score,greater_is_better=True,
                      needs_threshold=True)
    auc_scores = cross_val_score(log_reg, tfidf_w2v_std_train_data, y_train, cv=5, sc
    log_reg.fit(tfidf_w2v_std_train_data,y_train)
    y_pred_proba = log_reg.predict_proba(tfidf_w2v_std_train_data)[:,:1]
    train_auc = roc_auc_score(y_train, y_pred_proba)
    train_auc_values.append(train_auc)
    cv_scores.append(auc_scores.mean())

print ('train scores')
print (train_auc_values)
print (' '*50)
print ('CV scores')
print (cv_scores)

# changing to misclassification error

log = [math.log10(x) for x in alpha_values]
# plot misclassification error vs alpha
plt.plot(log, cv_scores,label='CV')
#plt.label('cv_f1')
plt.plot(log,train_auc_values,label='train')
#plt.label('train_f1')
plt.legend()
plt.xlabel('hyperparameter C')
plt.ylabel('AUC')
plt.show()

```


train scores

[0.8316270587156991, 0.843892458117977, 0.8469494345355281, 0.8522581075246595, 0.854040093152

\*\*\*\*\*

CV scores

[0.8265845605404006, 0.8399606573161972, 0.8429660551070757, 0.8476368060962354, 0.84920590096



```
In [110]: log_reg = LogisticRegression(C=0.01,penalty='l2',max_iter=10000,tol=0.001)
log_reg.fit(tfidf_w2v_std_train_data,y_train)
```

```
Out[110]: LogisticRegression(C=0.01, class_weight=None, dual=False, fit_intercept=True,
intercept_scaling=1, max_iter=10000, multi_class='ovr', n_jobs=1,
penalty='l2', random_state=None, solver='liblinear', tol=0.001,
verbose=0, warm_start=False)
```

```
In [111]: predictions = log_reg.predict(standardizing.transform(tfidf_w2v_test_data))
print('accuracy = {0}'.format(log_reg.score(standardizing.transform(tfidf_w2v_test_data),y_test)))
precision = precision_score(y_test, predictions,pos_label=1)
Recall = recall_score(y_test, predictions,pos_label=1)
f1 = f1_score(y_test, predictions,pos_label=1)
print ('precision = {0}'.format(precision))
print ('Recall={0}'.format(Recall))
print ('f1_score={0}'.format(f1))
```

```
accuracy = 76.42
```

```
precision = 0.8806308127780025
```

```
Recall=0.8408494208494208
```

```
f1_score=0.8602804661268022
```


```
In [112]: from sklearn.metrics import confusion_matrix
import seaborn as sns
```

```

result = confusion_matrix(y_test,predictions)
#print(result)
sns.set(font_scale=1.4)#for label size
sns.heatmap(result, annot=True,annot_kws={"size": 16}, fmt='g')

```


Out[112]: <matplotlib.axes.\_subplots.AxesSubplot at 0x449631beb8>



```

In [113]: y_pred_proba = log_reg.predict_proba(standardizing.transform(tfidf_w2v_test_data))[:
fpr, tpr, _ = roc_curve(y_test, y_pred_proba,pos_label=1 )
#auc = roc_auc_score(y_test, y_pred_proba)
auc = np.trapz(tpr,fpr)
plt.plot(fpr,tpr,label="data 1, auc="+str(auc))
plt.legend(loc=4)
plt.show()

```



## 10.2 [4.4.2] Applying Logistic Regression with L1 regularization on TFIDF W2V, SET 4

```
In [114]: from sklearn.metrics import make_scorer
          from sklearn.model_selection import cross_val_score
          from sklearn.metrics import roc_auc_score
          # creating list for hyperparameter alpha
          alpha_values = [0.0001,0.0005,0.001,0.005,0.01,0.05,0.1,100,500,1000,2500,5000,10000]
          # empty list that will hold cv scores

          cv_scores = []
          train_auc_values = []

          # perform 10-fold cross validation
          for alpha in alpha_values:
              log_reg = LogisticRegression(C=alpha,penalty='l1',max_iter=10000)
              auc = make_scorer(roc_auc_score,greater_is_better=True,
                               needs_threshold=True)
              auc_scores = cross_val_score(log_reg, tfidf_w2v_std_train_data, y_train, cv=5, s
              log_reg.fit(tfidf_w2v_std_train_data,y_train)
              y_pred_proba = log_reg.predict_proba(tfidf_w2v_std_train_data)[:,:1]
              train_auc = roc_auc_score(y_train, y_pred_proba)
              train_auc_values.append(train_auc)
              cv_scores.append(auc_scores.mean())
```

```

print ('train scores')
print (train_auc_values)
print (''*50)
print ('CV scores')
print (cv_scores)

# changing to misclassification error

log = [math.log10(x) for x in alpha_values]
# plot misclassification error vs alpha
plt.plot(log, cv_scores, label='CV')
#plt.label('cv_f1')
plt.plot(log, train_auc_values, label='train')
#plt.label('train_f1')
plt.legend()
plt.xlabel('hyperparameter C')
plt.ylabel('AUC')
plt.show()

```


train scores

[0.5633358424835658, 0.7540617737940192, 0.8097770102846455, 0.8395026884035292, 0.845292490702]

\*\*\*\*\*

CV scores

[0.5582377926129428, 0.7205708054481125, 0.7946056041449776, 0.8355777409967038, 0.84093418781]





```
In [115]: log_reg = LogisticRegression(C=0.05,penalty='l1',max_iter=10000,tol=0.001)
log_reg.fit(tfidf_w2v_std_train_data,y_train)
```


```
Out[115]: LogisticRegression(C=0.05, class_weight=None, dual=False, fit_intercept=True,
    intercept_scaling=1, max_iter=10000, multi_class='ovr', n_jobs=1,
    penalty='l1', random_state=None, solver='liblinear', tol=0.001,
    verbose=0, warm_start=False)
```

```
In [116]: predictions = log_reg.predict(standardizing.transform(tfidf_w2v_test_data))
print('accuracy = {0}'.format(log_reg.score(standardizing.transform(tfidf_w2v_test_d
precision = precision_score(y_test, predictions,pos_label=1)
Recall = recall_score(y_test, predictions,pos_label=1)
f1 = f1_score(y_test, predictions,pos_label=1)
print ('precision = {0}'.format(precision))
print ('Recall={0}'.format(Recall))
print ('f1_score={0}'.format(f1))
```


```
accuracy = 77.99333333333334
precision = 0.8777699475373894
Recall=0.8656370656370657
f1_score=0.8716612884413515
```

```
In [117]: from sklearn.metrics import confusion_matrix
import seaborn as sns
result = confusion_matrix(y_test,predictions)
#print(result)
sns.set(font_scale=1.4)#for label size
sns.heatmap(result, annot=True,annot_kws={"size": 16}, fmt='g')
```

```
Out[117]: <matplotlib.axes._subplots.AxesSubplot at 0x449d2f7a58>
```



```
In [118]: y_pred_proba = log_reg.predict_proba(standardizing.transform(tfidf_w2v_test_data))[:
fpr, tpr, _ = roc_curve(y_test, y_pred_proba, pos_label=1 )
#auc = roc_auc_score(y_test, y_pred_proba)
auc = np.trapz(tpr,fpr)
plt.plot(fpr,tpr,label="data 1, auc="+str(auc))
plt.legend(loc=4)
plt.show()
```



## 11 5) Tabular form of Results

In [119]: `from prettytable import PrettyTable`

```
x = PrettyTable()
```

```
x.field_names = ['Featurization','regularization','accuracy','AUC',
                  'precision','recall','f1_score']
x.add_row(['BOW','12','92.51','0.931','0.9667','0.8784','0.9202'])
x.add_row(['BOW','11','87.346','0.8829','0.8935','0.9688','0.9291'])
x.add_row(['TFIDF','12','90.47','0.9138','0.9223','0.9713','0.9462'])
x.add_row(['TF-IDF','11','91.13','0.937','0.916','0.9878','0.9505'])
x.add_row(['Avg W2V','12','81.023','0.7754','0.9091','0.86','0.8875'])
x.add_row(['Avg W2V','11','86.45','0.8557','0.86459','0.9974','0.9272'])
x.add_row(['TFIDF W2V','12','76.42','0.593','0.8808','0.8408','0.8602'])
x.add_row(['TFIDF W2V','11','77.99','0.590','0.8777','0.8656','0.8716'])
print (x)
```

Featurization	regularization	accuracy	AUC	precision	recall	f1_score
BOW	12	92.51	0.931	0.9667	0.8784	0.9202
BOW	11	87.346	0.8829	0.8935	0.9688	0.9291
TFIDF	12	90.47	0.9138	0.9223	0.9713	0.9462

TF-IDF	11	91.13	0.937	0.916	0.9878	0.9505	
Avg W2V	12	81.023	0.7754	0.9091	0.86	0.8875	
Avg W2V	11	86.45	0.8557	0.86459	0.9974	0.9272	
TFIDF W2V	12	76.42	0.593	0.8808	0.8408	0.8602	
TFIDF W2V	11	77.99	0.590	0.8777	0.8656	0.8716	
+-----+-----+-----+-----+-----+-----+-----+							

## 12 6) Conclusions

- 1) TFIDF with l1 regularizer Values has high values of preacision, recall , f1 score and it also has highest AUC.
- 2) BOW with l2 regularizer has auc close to TFIDF,but low recall and f1\_score and there is high multicollinearity in BOW with l2 regularizer
- 3) so, TFIDF with l1 regularizer is a best vector.